


CS-GY 6033

Design and Analysis of Algorithms I

Lecture 7 | Part 1

Shortest Paths



Recall

- ▶ The **length** of a path is

$(\# \text{ of nodes}) - 1$


Definitions

- ▶ A **shortest path** between u and v is a path between u and v with smallest possible length.
 - ▶ There may be several, or none at all.
- ▶ The **shortest path distance** is the length of a shortest path.
 - ▶ Convention: ∞ if no path exists.
 - ▶ “the distance between u and v ” means spd .

Today: Shortest Paths


- ▶ **Given:** directed/undirected graph G , source u
- ▶ **Goal:** find shortest path from u to every other node

Example



Key Property

- ▶ A shortest path of length k is composed of:
 - ▶ A **shortest path** of length $k - 1$.
 - ▶ Plus one edge.



Algorithm Idea

- ▶ Find all nodes distance 1 from source.
- ▶ Use these to find all nodes distance 2 from source.
- ▶ Use these to find all nodes distance 3 from source.
- ▶ ...

It turns out...

...this is exactly what BFS does.

CS-GY 6033

Design and Analysis of Algorithms I


Lecture 7 | Part 2

BFS for Shortest Paths

Key Property of BFS

- ▶ For any $k \geq 1$ you choose:
- ▶ All nodes distance $k - 1$ from source are added to the queue before any node of distance k .
- ▶ Therefore, nodes are “processed” (popped from queue) in order of distance from source.

Example



$P: [u]$

$P: [a, b]$ dist 1

$P: [\underbrace{b}_{\text{dist 1}}, \underbrace{c, d}_{\text{dist 2}}]$

Discovering Shortest Paths

- ▶ We “discover” shortest paths when we pop a node from queue and look at its neighbors.
- ▶ But the neighbor's status matters!

Consider This

- ▶ We pop a node s .
- ▶ It has a neighbor v whose status is **undiscovered**.
- ▶ We've discovered a **shortest path** to v through s !

Consider This

- ▶ We pop a node s .
- ▶ It has a neighbor v whose status is **pending** or **visited**.
- ▶ We already have a shortest path to v .

Modifying BFS

- ▶ Use BFS “framework”.
- ▶ Return dictionary of **search predecessors**.
 - ▶ If v is discovered while visiting u , we say that u is the **BFS predecessor** of v .
 - ▶ This encodes the shortest paths.
- ▶ Also return dictionary of shortest path distances.


```


def bfs_shortest_paths(graph, source):
    """Start a BFS at `source`."""
    status = {node: 'undiscovered' for node in graph.nodes}
    distance = {node: float('inf') for node in graph.nodes}
    predecessor = {node: None for node in graph.nodes}

    status[source] = 'pending'
    distance[source] = 0
    pending = deque([source])

    # while there are still pending nodes
    while pending:
        u = pending.popleft()
        for v in graph.neighbors(u):
            # explore edge (u,v)
            if status[v] == 'undiscovered':
                status[v] = 'pending'
                distance[v] = distance[u] + 1
                predecessor[v] = u
                # append to right
                pending.append(v)
        status[u] = 'visited'


    return predecessor, distance

```



$$dis[v] = dis[u] + 1$$

Example



$P: []$

$P: [u]$

$P: [a, b]$

$P: [b, c, d]$

$\vdots$

CS-GY 6033

Design and Analysis of Algorithms I

Lecture 7 | Part 3

BFS Trees

Result of BFS

- ▶ Each node reachable from source has a single BFS predecessor.
 - ▶ Except for the source itself.
- ▶ The result is a **tree** (or forest).


Trees

- ▶ A ~~(V, E)~~ **tree** is an undirected graph $T = (V, E)$ such that T is connected and $|E| = |V| - 1$.
- ▶ A **forest** is graph in which each connected component is a tree.




nodes = 1
edges = 0

Example

nodes = 6
edges = 5



Example



BFS Trees

- ▶ BFS trees and forests encode shortest path distances.

CS-GY 6033

Design and Analysis of Algorithms I


Lecture 7 | Part 4

Depth First Search


Visiting the Next Node

- ▶ Which node do we process next in a search?
- ▶ BFS: the **oldest** pending node.
- ▶ DFS (today): the **newest** pending node.
 - ▶ Naturally recursive.
 - ▶ Useful for solving different problems.

Example (BFS)




Example (DFS)



```
def dfs(graph, u, status=None):  
    """Start a DFS at `u`."""  
    # initialize status if it was not passed  
    if status is None:  
        status = {node: 'undiscovered' for node in graph.nodes}  
  
    status[u] = 'pending'  
    for v in graph.neighbors(u): # explore edge (u, v)  
        if status[v] == 'undiscovered':  
            dfs(graph, v, status)  
    status[u] = 'visited'
```

Exercise


Write the nested function calls for a DFS on the graph below.



```
def dfs(graph, u, status=None):  
    """Start a DFS at `u`."""  
    ...  
    status[u] = 'pending'  
    for v in graph.neighbors(u): # explore edge (u, v)  
        if status[v] == 'undiscovered':  
            dfs(graph, v, status)  
    status[u] = 'visited'
```

Exercise

Write the nested function calls for a DFS on the graph below.



DFS (a):

DFS (b):

DFS (d):

DFS(e):

DFS(g)

DFS(h)

DFS(i)

DFS(c):
DFS(f)

Differences

- ▶ In **BFS**, we “finish” a node u before moving on to the next.
- ▶ In **DFS**, we go to many other nodes, but “come back” to u .

Main Idea

We'll see that the nested structure of the **recursive function calls** gives us useful new information about the graph's structure.

Full DFS



- ▶ `dfs(u)` will visit all nodes **reachable** from u .
 - ▶ But not all nodes may be reachable from u !
- ▶ To visit **all** nodes in graph, need **full DFS**.

```
def full_dfs(graph):  
    status = {node: 'undiscovered' for node in graph.nodes}  
    for node in graph.nodes:  
        if status[node] == 'undiscovered'  
            dfs(graph, node, status)
```

```
def full_dfs(graph):
    status = {node: 'undiscovered' for node in graph.nodes}
    for node in graph.nodes:
        if status[node] == 'undiscovered':
            dfs(graph, node, status)

def dfs(graph, u, status=None):
    """Start a DFS at `u`."""
    # initialize status if it was not passed
    if status is None:
        status = {node: 'undiscovered' for node in graph.nodes}

    status[u] = 'pending'
    for v in graph.neighbors(u): # explore edge (u, v)
        if status[v] == 'undiscovered':
            dfs(graph, v, status)
    status[u] = 'visited'
```



Time Complexity

- ▶ In a full DFS:
 - ▶ dfs called on each node exactly once.
 - ▶ Like BFS, each edge is explored exactly:
 - ▶ once if directed
 - ▶ twice if undirected
- ▶ Time: $\Theta(V + E)$, **just like BFS.**

CS-GY 6033


Design and Analysis of Algorithms I

Lecture 7 | Part 5

Nesting Properties of DFS


Exercise

True or False: if v is reachable from u and v is **undiscovered** when $\text{dfs}(u)$ is called, then $\text{dfs}(v)$ must be called during $\text{dfs}(u)$.




False!

- ▶ Suppose $\text{dfs}(4)$ is the root call.
 - ▶ When $\text{dfs}(1)$ is called, node 5 is undiscovered.
 - ▶ But $\text{dfs}(5)$ is **not** called during $\text{dfs}(1)$.



However..

- This intuition is correct if there is a path of **undiscovered** nodes from u to v when $\text{dfs}(u)$ is called.



Key Property of DFS (Informal)

- ▶ If at the time $\text{dfs}(u)$ is called...
 1. v is **undiscovered**; and
 2. there is a path of **undiscovered** nodes from u to v ,
- ▶ ...then $\text{dfs}(v)$ will **start and finish** during the call to $\text{dfs}(u)$.

Start and Finish Times

- ▶ Keep a running clock (an integer).
- ▶ For each node, record
 - ▶ **Start time**: time when marked **pending**
 - ▶ **Finish time**: time when marked **visited**
- ▶ Increment clock whenever node is marked **pending/visited**


```
from dataclasses import dataclass

@dataclass
class Times:
    clock: int
    start: dict
    finish: dict


def full_dfs_times(graph):
    status = {node: 'undiscovered' for node in graph.nodes}
    predecessor = {node: None for node in graph.nodes}
    times = Times(clock=0, start={}, finish={})
    for u in graph.nodes:
        if status[u] == 'undiscovered':
            dfs_times(graph, u, status, times)
    return times, predecessor

def dfs_times(graph, u, status, predecessor, times):
    times.clock += 1
    times.start[u] = times.clock
    status[u] = 'pending'
    for v in graph.neighbors(u): # explore edge (u, v)
        if status[v] == 'undiscovered':
            predecessor[v] = u
            dfs_times(graph, v, status, times)
    status[u] = 'visited'
    times.clock += 1
    times.finish[u] = times.clock
```

Example



Key Property of DFS




- Suppose $\text{dfs}(u)$ is called before $\text{dfs}(v)$.
- If when $\text{dfs}(u)$ is called there is a path of **undiscovered** nodes from u to v , then:

$$\text{start}[u] < \text{start}[v] < \text{finish}[v] < \text{finish}[u].$$

- Otherwise:

$$\text{start}[u] < \text{finish}[u] < \text{start}[v] < \text{finish}[v].$$



Key Property

- ▶ Take any two nodes u and v ($u \neq v$).
- ▶ Assume for simplicity that $\text{start}[u] < \text{start}[v]$.
- ▶ Exactly one of these is true:
 - ▶ $\text{start}[u] < \text{start}[v] < \text{finish}[v] < \text{finish}[u]$
 - ▶ $\text{start}[u] < \text{finish}[u] < \text{start}[v] < \text{finish}[v]$

CS-GY 6033


Design and Analysis of Algorithms I

Lecture 7 | Part 6

Cycles


Cycle

- A **cycle** in a directed graph is a path that starts and ends at the same node.




Cycle

- Alternatively: there is a **cycle** if u is reachable from v and v is reachable from u , for some $u \neq v$.




DAG

- A **directed acyclic graph** (DAG) is a directed graph with **no cycles**.



Cyclic Graphs

- ▶ A graph is cyclic even if it has only one cycle.
 - ▶ It doesn't have to be the whole graph.




Detecting Cycles

- ▶ We check for cycles by looking for **back edges** in a full DFS.
- ▶ (u, v) is a **back edge** if while visiting node u , we see that v is **pending**.

```
...  
for v in graph.neighbors(u): # explore edge (u, v)  
    if status[v] == 'undiscovered':  
        dfs(graph, v, status)  
    elif status[v] == 'pending':  
        # back edge (u, v) found!  
...  

```

Example



Theorem

*A directed graph has a cycle **if (and only if)** a full DFS finds a back edge.*

Why?

- ▶ If a back edge (u, v) is found, then a cycle exists.
 - ▶ Suppose v is pending when we visit u .
 - ▶ This means that there is a path from v to u .
 - ▶ There is also a path from u to v .
 - ▶ So there is a cycle.

Why?

- ▶ If a cycle exists, then there is a back edge.
 - ▶ Suppose there is a cycle $v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_k \rightarrow v_1$.
 - ▶ Without loss of generality, assume v_1 is the first node in the cycle that is visited by the full DFS.
 - ▶ At the moment of $\text{dfs}(v_1)$, there is a path of undiscovered nodes between v_1 and v_k .
 - ▶ Therefore $\text{dfs}(v_k)$ will be called during $\text{dfs}(v_1)$.
 - ▶ During $\text{dfs}(v_k)$, we'll see the back edge.

CS-GY 6033


Design and Analysis of Algorithms I

Lecture 7 | Part 7


Topological Sort

Applications of DFS

- ▶ Is node v reachable from node u ? **DFS, BFS**
- ▶ Is the graph connected? **DFS, BFS**
- ▶ How many connected components? **DFS, BFS**
- ▶ Find the shortest path between u and v . **DFS, BFS**
- ▶ Does the graph have a cycle? **DFS, BFS**



Prerequisite Graphs



Goal: find order in which classes should be taken in order to satisfy the prerequisites of CSCI 480.

Note

- ▶ Prerequisite graphs are¹ DAGs.

¹Or they should be, at least!


Topological Sorts

► **Given:** a DAG, $G = (V, E)$.

→ ► **Compute:** an ordering of V such that if $(u, v) \in E$, then u comes before v in the ordering

► This is called a **topological sort** of G .


Example



✓ MATH 121, CSCI 101, CSCI 202, CSCI 102, CSCI 303, CSCI 201, CSCI 310, CSCI 480

Computing a Topological Sort

- ▶ How do we compute a topological sort, algorithmically?
- ▶ **Observation:** if v is reachable from u , v **must** come **after** u in the topological sort.



Exercise

Suppose v is reachable from u in a DAG. ✓

True or false: after a full DFS, $\text{finish}[v] < \text{finish}[u]$.

Claim

- 
- If v is reachable from u in a DAG, then:

$$\text{finish}[v] < \text{finish}[u]$$


Idea


- ▶ Take any two nodes u and v ($u \neq v$).
- ▶ Assume the graph is a DAG, run DFS.
- ▶ If v is reachable from u , then $\text{finish}[v]$ $<$ $\text{finish}[u]$.

Putting it together...

- ▶ **Observation:** If v is reachable from u , then v must come after u in the topological sort.
- ▶ **Recall:** If v is reachable from u , then $\text{finish}[v] < \text{finish}[u]$.

Exercise

Compute start and finish times using CSCI 101 as the source.



Idea

- ▶ **Observation:** If v is reachable from u , then v must come after u in the topological sort.
- ▶ **Recall:** If v is reachable from u , then $\text{finish}[v] < \text{finish}[u]$.
- ▶ **Therefore:** if $\text{finish}[v] < \text{finish}[u]$, then v must come after u in the topological sort.
- ✓▶ **Idea:** sort nodes in **descending** order by finish time.


Algorithm

- ▶ To find a topological sort (if it exists):
 - ▶ Compute times with Full DFS. $\Theta(|V| + |E|)$
 - ▶ Sort in **descending** order by finish time. $\Theta(|V|)$
- ▶ Time complexity:

$$\Theta(|V| + |E|)$$

8/29/8

Example



Note

- ▶ There can be many valid topological sorts!